

Supplemental Materials for “Another Vertical View: A Hierarchical Network for Heterogeneous Trajectory Prediction via Spectrums”

1 TRANSFORMER DETAILS

We employ the Transformer [1] as the backbone to encode trajectory spectrums and the scene context in the two proposed sub-networks. The Transformer used in the E-V²-Net has two main parts, the Transformer Encoder and the Transformer Decoder, both of which are made up of several attention layers.

Attention Layers. Multi-Head Attention operations are applied in each of the attention layers. Following [1], each layer’s multi-head dot product attention with H heads is calculated as:

$$\text{Attention}(\mathbf{q}, \mathbf{k}, \mathbf{v}) = \text{softmax}\left(\frac{\mathbf{q}\mathbf{k}^T}{\sqrt{d}}\right) \mathbf{v}, \quad (1)$$

$$\text{MultiHead}(\mathbf{q}, \mathbf{k}, \mathbf{v}) = \text{fc}\left(\text{concat}(\{\text{Attention}_i(\mathbf{q}, \mathbf{k}, \mathbf{v})\}_{i=1}^H)\right). \quad (2)$$

Here, $\text{fc}()$ denotes one fully connected layer that concatenates all heads’ outputs. Query matrix \mathbf{q} , key matrix \mathbf{k} , and value matrix \mathbf{v} , are the three layer inputs. Each attention layer also contains an MLP (denoted as MLP_a) to extract the attention features further. It contains two fully connected layers. ReLU activations are applied in the first layer. Formally, we have output feature \mathbf{f}_o of this layer:

$$\mathbf{f}_o = \text{ATT}(\mathbf{q}, \mathbf{k}, \mathbf{v}) = \text{MLP}_a(\text{MultiHead}(\mathbf{q}, \mathbf{k}, \mathbf{v})). \quad (3)$$

Transformer Encoder. The transformer encoder comprises several encoder layers, and each encoder layer contains an attention layer and an encoder MLP (MLP_e). Residual connections and normalization operations are applied to prevent the network from overfitting. Let $\mathbf{h}^{(l+1)}$ denote the output of l -th encoder layer, and $\mathbf{h}^{(0)}$ denote the encoder’s initial input. For l -th encoder layer, the calculation of the layer output $\mathbf{h}^{(l+1)}$ can be written as:

$$\begin{aligned} \mathbf{a}^{(l)} &= \text{ATT}(\mathbf{h}^{(l)}, \mathbf{h}^{(l)}, \mathbf{h}^{(l)}) + \mathbf{h}^{(l)}, \\ \mathbf{a}_n^{(l)} &= \text{Normalization}(\mathbf{a}^{(l)}), \\ \mathbf{c}^{(l)} &= \text{MLP}_e(\mathbf{a}_n^{(l)}) + \mathbf{a}_n^{(l)}, \\ \mathbf{h}^{(l+1)} &= \text{Normalization}(\mathbf{c}^{(l)}). \end{aligned} \quad (4)$$

Transformer Decoder. Similar to the Transformer encoder, the Transformer decoder comprises several decoder layers,

and each is stacked with two different attention layers. The first attention layer in the Transformer decoder focuses on the essential parts in the Transformer encoder’s outputs \mathbf{h}_e queried by the decoder’s input \mathbf{X} . The second layer is the same self-attention layer as in the encoder. Similar to Eq. 4, we have the decoder layer’s output feature $\mathbf{h}^{(l+1)}$:

$$\begin{aligned} \mathbf{a}^{(l)} &= \text{ATT}(\mathbf{h}^{(l)}, \mathbf{h}^{(l)}, \mathbf{h}^{(l)}) + \mathbf{h}^{(l)}, \\ \mathbf{a}_n^{(l)} &= \text{Normalization}(\mathbf{a}^{(l)}), \\ \mathbf{a}_2^{(l)} &= \text{ATT}(\mathbf{h}_e, \mathbf{h}^{(l)}, \mathbf{h}^{(l)}) + \mathbf{h}^{(l)}, \\ \mathbf{a}_{2n}^{(l)} &= \text{Normalization}(\mathbf{a}_2^{(l)}) \\ \mathbf{c}^{(l)} &= \text{MLP}_d(\mathbf{a}_{2n}^{(l)}) + \mathbf{a}_{2n}^{(l)}, \\ \mathbf{h}^{(l+1)} &= \text{Normalization}(\mathbf{c}^{(l)}). \end{aligned} \quad (5)$$

Positional Encoding. Before feeding agents representations or trajectory spectrums into the Transformer, we add the positional coding to inform the relative position of each timestep or frequency portion in the sequential inputs. The position coding \mathbf{f}_e^t at step t ($1 \leq t \leq t_h$) is obtained by:

$$\begin{aligned} \mathbf{f}_e^t &= (f_{e0}^t, \dots, f_{ei}^t, \dots, f_{ed-1}^t) \in \mathbb{R}^d, \\ \text{where } f_{ei}^t &= \begin{cases} \sin(t/10000^{d/i}), & i \text{ is even;} \\ \cos(t/10000^{d/(i-1)}), & i \text{ is odd.} \end{cases} \end{aligned} \quad (6)$$

Then, we have the positional coding matrix \mathbf{f}_e that describes t_h steps of sequences:

$$\mathbf{f}_e = (\mathbf{f}_e^1, \mathbf{f}_e^2, \dots, \mathbf{f}_e^{t_h})^T \in \mathbb{R}^{t_h \times d}. \quad (7)$$

The final Transformer input \mathbf{X}_T is the addition of the original sequential input \mathbf{X} and the positional coding matrix \mathbf{f}_e . Formally,

$$\mathbf{X}_T = \mathbf{X} + \mathbf{f}_e \in \mathbb{R}^{t_h \times d}. \quad (8)$$

Layer Configurations. We employ $L = 4$ layers of encoder-decoder structure with $H = 8$ attention heads in each Transformer-based sub-networks. The MLP_e and the MLP_d have the same shape. Both of them consist of two fully connected layers. The first layer has 512 output units with the ReLU activation, and the second layer has 128 but does not use any activations. The output dimensions of fully connected layers used in multi-head attention layers are set to $d = 128$.




Fig. 1. Matrices views of a trajectory (2D bounding box) and an image.

2 LINEAR LEAST SQUARES TRAJECTORY PREDICTION

The linear least squares trajectory prediction method aims to minimize the mean square error between the predicted and agents' groundtruth trajectories. When predicting, we perform a separate least squares operation for each dimension of the M -dimensional observed trajectory \mathbf{X} . Simply, we want to find the $\mathbf{x}_m = (b_m, w_m)^T \in \mathbb{R}^2$ ($1 \leq m \leq M$), such that

$$\hat{\mathbf{Y}} = (\hat{\mathbf{Y}}_1, \hat{\mathbf{Y}}_2, \dots, \hat{\mathbf{Y}}_m, \dots, \hat{\mathbf{Y}}_M) \in \mathbb{R}^{t_f \times M},$$

$$\text{where } \hat{\mathbf{Y}}_m = \mathbf{A}_f \mathbf{x}_m = \begin{pmatrix} 1 & t_h + 1 \\ 1 & t_h + 2 \\ \dots & \dots \\ 1 & t_h + t_f \end{pmatrix} \begin{pmatrix} b_m \\ w_m \end{pmatrix}. \quad (9)$$

For one of agents' observed M -dimensional trajectory $\mathbf{X} \in \mathbb{R}^{t_h \times M}$, we have the trajectory slice on the m -th dimension

$$\mathbf{X}_m = (r_{m1}, r_{m2}, \dots, r_{mt_h})^T. \quad (10)$$

Suppose we have a coefficient matrix \mathbf{A}_h , where

$$\mathbf{A}_h = \begin{pmatrix} 1 & 1 & 1 & \dots & 1 \\ 1 & 2 & 3 & \dots & t_h \end{pmatrix}^T. \quad (11)$$

We aim to find a $\mathbf{x}_m \in \mathbb{R}^2$, such that the mean square $\|\mathbf{A}_h \mathbf{x}_m - \mathbf{X}_m\|_2^2$ could reach its minimum value. Under this condition, we have

$$\mathbf{x}_m = (\mathbf{A}_h^T \mathbf{A}_h)^{-1} \mathbf{A}_h^T \mathbf{X}_m. \quad (12)$$

Then, we have the predicted m -th dimension trajectory

$$\hat{\mathbf{Y}}_m = \mathbf{A}_f \mathbf{x}_m. \quad (13)$$

The final M -dimensional predicted trajectory $\hat{\mathbf{Y}}$ is obtained by stacking all results. Formally,

$$\hat{\mathbf{Y}} = \mathbf{A}_f (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_M). \quad (14)$$

3 2D DFT V.S. BILINEAR STRUCTURE

We apply different transforms on *each dimension* of the trajectory to obtain the corresponding trajectory spectrums either in V²-Net or the enhanced E-V²-Net. Moreover, considering that one of our main contributions is to establish connections between trajectories (or spectrums) of different dimensions, a more natural idea might be to apply some 2D transform directly to these trajectories. However, it appears to be less effective from both theoretical analyses and experimental results. In this section, we will discuss the discrepancy between the 2D transform and the proposed bilinear structure in describing the two factors, including

the frequency response of the trajectory and the dimension-wise interactions, from different perspectives, taking DFT as an example.

DFT on Different Directions in Trajectories. The 2D DFT can be decomposed into two consecutive 1D DFTs performed in different directions of the target 2D matrix. The M -dimensional trajectory $\mathbf{X} \in \mathbb{R}^{N \times M}$ is also a 2D matrix similar to 2D grayscale images. Although the 2D DFT and its variations have achieved impressive results in tasks related to image processing, they might be not directly applied to trajectories. We will analyze this problem specifically by focusing on the different directions of the transforms in the trajectory.

Fig. 1 shows an $M = 4$ 2D bounding box trajectory and an image with the matrix view. As shown in Fig. 1 (b), whether the image is sliced horizontally or vertically, the resulting vector could reflect the change in grayscale values in a particular direction. Therefore, when performing the 2D transform, the first 1D transform will extract the frequency response in a specific direction, while the second 1D transform will fuse it with the frequency response in the vertical direction.

In contrast, different slice directions of the trajectory may lead to different meanings. If the trajectories are sliced according to the time dimension, then four 1D time series will be obtained as shown in Fig. 1 (a). Applying 1D transforms to these four sequences, we can obtain four trajectory spectrums that could describe agents' frequency responses and thus describe their motions from the global plannings and interaction details at different scales. However, if the trajectory is sliced from the dimensional direction, then N (6 in the figure) 4-dimensional vectors will be obtained. These vectors contain information about agents' locations and postures at a particular moment. In addition, the focused dimension-wise interactions are also contained in these vectors. However, it should be noted that we are more interested in the relationships between the data in these vectors, *i.e.*, the "edges" between the different data. If a 1D transform is applied to these vectors, the resulting spectrum may hardly have a clear physical meaning, because the temporal or spatial adjacencies of these points are not reflected in these 4-dimensional vectors.

For example, suppose we want to apply the 1D DFT on the 4-dimensional (2D bounding box) vector $\mathbf{x} = (x_l, y_l, x_r, y_r)^T$. Simply, we have:

$$\begin{aligned} \mathcal{X} = \text{DFT}[\mathbf{x}] &= \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & -j & -1 & j \\ 1 & -1 & 1 & -1 \\ 1 & j & -1 & -j \end{pmatrix} \begin{pmatrix} x_l \\ y_l \\ x_r \\ y_r \end{pmatrix} \\ &= \begin{pmatrix} x_l + y_l + x_r + y_r \\ x_l - x_r - j(y_l - y_r) \\ x_l - y_l + x_r - y_r \\ x_l - x_r + j(y_l - y_r) \end{pmatrix}. \end{aligned} \quad (15)$$

Accordingly, we have its fundamental frequency portion $\mathcal{X}[0] = x_l + y_l + x_r + y_r$ and the high-frequency portion $\mathcal{X}[2] = x_l - y_l + x_r - y_r$. However, since the four positions $\{x_l, y_l, x_r, y_r\}$ do not have specific time-dependent or space-dependent like time-sequences and images, these




Fig. 2. Visualized comparisons of 2D DFT bilinear structure.

frequency components may hardly reflect the specific frequency response. For example, the fundamental frequencies can represent their average value for a time series, yet the values obtained by directly summing the 4 position coordinates of the 2 points of the 2D bounding box would be uninterpretable. In other words, each element in this 4-dimensional vector is relatively independent, and their connection relationships are more like a *graph* rather than a sequence where an order is required.

Quantitative Analyses. To verify our thoughts, we perform ablation experiments on SDD and nuScenes to compare the effects of 2D DFT and the bilinear structure quantitatively. As shown in TABLE 1, the results of APP2 and APP3 (or APP5 and APP6) show that 2D DFT does not improve quantitative trajectory prediction performance as effectively as bilinear structures. On the contrary, in the prediction of the more complex 3D bounding boxes ($M = 6$), using 2D DFT instead degrades the prediction performance compared to 1D DFT when no bilinear structures are used. These experimental results validate our thoughts of not using 2D transforms but bilinear structures.

Qualitative Analyses. We visualize the prediction results of different models under the effect of 2D DFT and bilinear structure qualitatively. As shown in Fig. 2, the V²-Net (2D DFT) performs not as well as E-V²-Net in both the prediction of agent motions and the interactions within the bounding box. In detail, predictions given by V²-Net (2D DFT) capture fewer path possibilities in the top left prediction scenario. In addition, some predicted trajectories are with less smoothness and naturalness. For example, the prediction in color #DF6091 to the left of the bottom left prediction scene gives a turn with a large angle to observation, which could be not physical-acceptable in the actual scenario. In contrast, predictions given by E-V²-Net have not shown similar results in this scenario. On the other hand, as shown in the traffic circle prediction scenario on the right, the shape of the bounding box is not well maintained in V²-Net’s predictions, such as the prediction in color #93C3CA to turn right to across the street.

4 A GRAPH-VIEW OF BILINEAR STRUCTURE

As mentioned above, the modeling of dimension-wise interactions focuses more on the relations of different trajectory dimensions, which could be difficult to represent by the 1D transform. The bilinear structure used in this manuscript

learns this connection relation through the outer product, pooling, and fully connected networks. To make it easier to understand, we further explain it from a graph view.

Given an undirected graph $G(t) = (V(t), E(t))$ with a time variable t , where $V(t)$ is the set of vertices, which contains all the position information of one agent at time t , and $E(t)$ is the set of edges, which represents the connection relationships between these vertices at time t . Formally,

$$\begin{aligned} V(t) &= \{f(r_{1t}), f(r_{2t}), \dots, f(r_{Mt})\} \\ &= \{f_{1,t}, f_{2,t}, \dots, f_{M,t}\}. \end{aligned} \quad (16)$$

where

$$f_{m,t} = f(r_{mt}) \in \mathbb{R}^d \quad (17)$$

indicates an embedding function to map these vertices into the high-dimension feature space. To establish and learn the connections between these vertices, we define the trainable adjacency matrix as:

$$A(t) = \begin{pmatrix} W_{1,1}(t) & \cdots & W_{1,M}(t) \\ \vdots & \ddots & \vdots \\ W_{M,1}(t) & \cdots & W_{M,M}(t) \end{pmatrix}. \quad (18)$$

Here, each matrix $W_{i,j} \in \mathbb{R}^{d \times d}$ are made trainable. They are used to describe the relation between node i and node j (i.e., $f_{i,t}$ and $f_{j,t}$).

We converge the information on the node edges by graph convolution. Formally,

$$f'_{m,t} = \sigma \left(W'_m \text{Flatten} \left(\sum_{j=1}^M W_{m,j}(t) f_{m,t} \otimes f_{j,t} \right) \right), \quad (19)$$

where σ represents a non-linear activation, and the W'_m is another trainable weight matrix. Finally, we have the refined vertices

$$V'(t) = \{f'_{1,t}, f'_{2,t}, \dots, f'_{M,t}\}. \quad (20)$$

It is worth noting that the above Eq. 19 and the bilinear structure introduced in the manuscript describe the same network structure, despite their difference in representation. To reduce unnecessary misunderstandings, we do not describe the network inference process through this graph form in the manuscript, although the use of a graph may make it easier to understand the motivation for the use of the bilinear model. In addition, all the operations above are performed on time series. If we use trajectory spectrums to replace the Eq. 17, and take the frequency variable n to instead the time variable t , we have

$$V(n) = \{f_{m,n}\}_{m=1}^M, \quad \text{where } f_{m,n} = f(s_{n,m}). \quad (21)$$

Accordingly, we have the refined vertices’ spectrum representations:

$$V'(t) = \{f'_{1,n}, f'_{2,n}, \dots, f'_{M,n}\}. \quad (22)$$

Then, the outer product matrix $R[n, :, :]$ has become the adjacency matrix of the graph $G(n) = (V(n), E(n))$ on the frequency node $n \in [1, N_h]$.

It is worth noting that the methods proposed in the manuscript do not really use graph structures and graph convolution operations. The analyses using the graph views are only intended to make it easier to understand the motivation and the rough working of bilinear structures. Therefore, this part of the analysis is for reference only.

TABLE 1
Validation of 2D DFT and bilinear structures with *best-of-20* on SDD (2D bounding box) the nuScenes (3D bounding box).

No.	Model	Type	T	N_{key}	BS	Dataset	ADE ↓	FDE ↓	AIoU ↑	FIoU ↑
APP1	V ² -Net	bb	DFT	3	×	SDD (2D bounding box)	6.78	10.73	0.717	0.601
APP2	V ² -Net	bb	2D DFT	3	×		6.74	10.84	0.723	0.602
APP3	E-V ² -Net	bb	DFT	3	✓		6.62	10.57	0.725	0.604
APP4	V ² -Net	3dbb	DFT	2	×	nuScenes (3D bounding box)	0.229	0.335	0.747	0.666
APP5	V ² -Net	3dbb	2D DFT	2	×		0.234	0.341	0.739	0.656
APP6	E-V ² -Net	3dbb	DFT	2	✓		0.210	0.300	0.762	0.688

TABLE 2
Comparisons of 3D skeleton prediction performance on Human3.6M.
Reported metrics are the FDE during different prediction periods
(length of these periods are shown in million seconds).

Models	Source	FDE ↓			
		80	160	320	400
Res. sup. [2]	CVPR 2017	34.7	62.0	101.1	115.5
Traj-GCN [3]	ICCV 2019	26.1	52.3	63.5	63.5
DMGNN [4]	CVPR 2020	33.6	65.9	79.7	79.7
MSR-GCN [5]	ICCV 2021	12.1	25.6	51.6	62.9
PGBIG [6]	CVPR 2022	10.3	22.7	47.4	58.5
SPGSN [7]	ECCV 2022	10.4	22.3	47.1	58.3
EqMotion [8]	CVPR 2023	9.1	20.1	43.7	55.0
E-V ² -Net-Haar	(Ours)	13.0	25.5	50.1	62.8

5 EVOLUTION OF MOTION (SKELETON) PREDICTION

The E-V²-Net is proposed to handle both trajectories' frequency response and the dimension-wise interactions. We have validated the proposed models with three different forms of trajectories, including **2D coordinates** ($M = 2$), **2D bounding boxes** ($M = 4$), and **3D bounding boxes** ($M = 6$). In order to verify more directly the modeling and prediction capabilities of the proposed model for heterogeneous trajectories, we validated the prediction performance for 3D human skeletons in this section. The skeletons we used for validation consist of 17 3D points, which we call **3D skeleton-17** ($M = 51$). It should be noted that the proposed models (V²-Net and E-V²-Net) are not specifically designed for motion (skeleton) prediction. We just want to verify our idea under a new type of heterogeneous trajectories.

Datasets. Following recent motion prediction approaches like [8], we choose the **Human3.6M** [9], [10] dataset to validate the motion prediction performance. It is a big dataset with 3.6 million 3D human poses and corresponding images, performed by 11 professional actors in 17 scenarios (such as discussion, smoking, taking photos, and talking on the phone). Its videos and pose data are recorded at 50Hz.

Baselines. We choose the following methods as baselines to validate E-V²-Net's skeleton prediction performance, including Res. sup. [2], Traj-GCN [3], DMGNN [4], MSR-GCN [5], PGBIG [6], SPGSN [7], EqMotion [8].

Metrics. We use the FDE (Final Displacement Error) as the metric to measure the 3D skeleton prediction performance. It should be noted that in the field of motion prediction, this metric is also more commonly known as the Mean Per Joint Position Error (MPJPE). For easier understanding, we still use the FDE as usual for trajectory prediction.

Implementation Details. Following previous settings, we use all data from subjects $\{1, 6, 7, 8, 9\}$ to train the model, subjects $\{11\}$ to validate, and subjects $\{5\}$ to test. When making the training samples, we sample observations with the frequency of 25Hz (*i.e.*, the sample interval is 40ms) and use $t_h = 10$ frames (400ms) of observations from all subjects to predict their possible trajectories (3D skeleton-17) for the next $t_f = 10$ frames (400ms).

We set $N_{key} = 4$, and $\{t_1^{key}, t_2^{key}, t_3^{key}, t_4^{key}\} = \{t_h + 1, t_h + 4, t_h + 7, t_h + 10\}$. The input dimension of the network is set to $M = 17 * 3 = 51$. We extend the feature dimension from 128 to 512 for each layer in the network to expand the model capacity. We disable the noise sampling layers so that the network could predict the deterministic predictions. In addition, since there is only one subject in the scene, all modules of social interaction and scene interaction are also disabled. When training on Human3.6M, we set the learning rate to 0.0005 and train the model for 200 epochs.

Comparisons to State-of-the-Art Methods. We show the comparisons of the proposed E-V²-Net Haar and several state-of-the-art motion prediction baselines in TABLE 2. The proposed model and MSR-GCN have similar motion prediction performance. Although the proposed E-V²-Net does not outperform the recently proposed methods like EqMotion (for about 14% performance drop), it still shows a strong competitive performance.

Qualitative Analysis. We show the visualized 3D skeleton prediction results in Fig. 3 to demonstrate how the proposed model handles the complex dimension-wise interactions within the trajectory. Naturally, dimension-wise interactions in a skeleton manifest as changes and interactions of edges between different joints, and are limited by the physical constraints and motions of the human body. The proposed model, although not purposely designed for motion prediction, still exhibits amazing prediction results. As shown in Fig. 3 (a), E-V²-Net successfully predicted the subsequent movements of the person who was running. It is particularly noteworthy that it has a better prediction of the legs in the skeleton, which is also shown in Fig. 3 (b) and (c). Unfortunately, the model does not predict the human arms very well, as shown in Fig. 3 (c) they remain almost stationary without any motion.

Summary. The prediction results in the more complex human 3D skeletons ($M = 51$) also demonstrate the effectiveness of the proposed E-V²-Net in dealing with dimension-wise interaction in more complex heterogeneous trajectories.

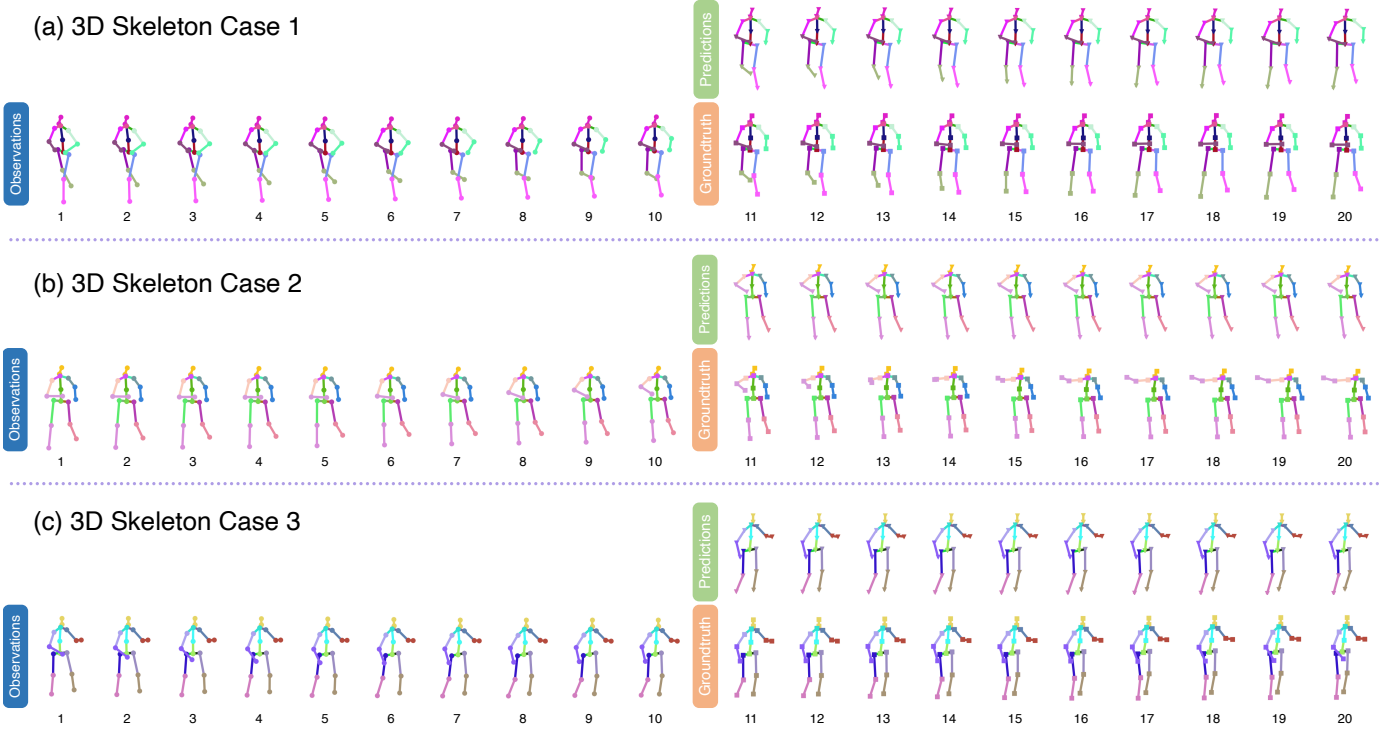


Fig. 3. Visualized 3D skeletons predictions (3D skeleton-17) on Human3.6M.

It also shows the higher trajectory prediction potential of the model without changing the model structure, which further demonstrates the generality of the proposed model “from another view”. However, it is worth noting that motion prediction is currently a challenging task, which is also more different from human trajectory prediction in terms of concerns and applications. While comparisons across tasks may be inappropriate, we only try to validate the model’s ability to handle dimension-wise interactions.

REFERENCES

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in neural information processing systems*, 2017, pp. 5998–6008. 1
- [2] J. Martinez, M. J. Black, and J. Romero, “On human motion prediction using recurrent neural networks,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2017, pp. 2891–2900. 4
- [3] W. Mao, M. Liu, M. Salzmann, and H. Li, “Learning trajectory dependencies for human motion prediction,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019, pp. 9489–9497. 4
- [4] M. Li, S. Chen, Y. Zhao, Y. Zhang, Y. Wang, and Q. Tian, “Dynamic multiscale graph neural networks for 3d skeleton based human motion prediction,” in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2020, pp. 214–223. 4
- [5] L. Dang, Y. Nie, C. Long, Q. Zhang, and G. Li, “Msr-gcn: Multi-scale residual graph convolution networks for human motion prediction,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021, pp. 11 467–11 476. 4
- [6] T. Ma, Y. Nie, C. Long, Q. Zhang, and G. Li, “Progressively generating better initial guesses towards next stages for high-quality human motion prediction,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022, pp. 6437–6446. 4
- [7] M. Li, S. Chen, Z. Zhang, L. Xie, Q. Tian, and Y. Zhang, “Skeleton-parted graph scattering networks for 3d human motion prediction,” in *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part VI*. Springer, 2022, pp. 18–36. 4
- [8] C. Xu, R. T. Tan, Y. Tan, S. Chen, Y. G. Wang, X. Wang, and Y. Wang, “Eqmotion: Equivariant multi-agent motion prediction with invariant interaction reasoning,” *arXiv preprint arXiv:2303.10876*, 2023. 4
- [9] C. Ionescu, D. Papava, V. Olaru, and C. Sminchisescu, “Human3.6m: Large scale datasets and predictive methods for 3d human sensing in natural environments,” *IEEE transactions on pattern analysis and machine intelligence*, vol. 36, no. 7, pp. 1325–1339, 2013. 4
- [10] C. S. Catalin Ionescu, Fuxin Li, “Latent structured models for human pose estimation,” in *International Conference on Computer Vision*, 2011. 4